

PARTE 3 - Configuración en Cluster Hadoop de Conexión con ES

“Captura de pantalla del proceso de configuración en Cluster Hadoop de Conexión con ES completo”

PART 1-



SSH en el navegador

↑ SUBIR ARCHIVO

↓ DESCARGAR ARCHIVO



Welcome to Ubuntu 22.04.5 LTS (GNU/Linux 6.8.0-1025-gcp x86_64)

* Documentation: <https://help.ubuntu.com>
* Management: <https://landscape.canonical.com>
* Support: <https://ubuntu.com/pro>

System information as of Sun Mar 23 18:53:34 UTC 2025

System load:	0.25	Processes:	137
Usage of /:	31.0% of 48.27GB	Users logged in:	1
Memory usage:	67%	IPv4 address for ens3:	10.132.0.6
Swap usage:	0%		

=> There is 1 zombie process.

* Strictly confined Kubernetes makes edge and IoT secure. Learn how MicroK8s just raised the bar for easy, resilient and secure K8s cluster deployment.

<https://ubuntu.com/engage/secure-kubernetes-at-the-edge>

Expanded Security Maintenance for Applications is not enabled.

0 updates can be applied immediately.

Enable ESM Apps to receive additional future security updates.
See <https://ubuntu.com/esm> or run: `sudo pro status`

New release '24.04.2 LTS' available.
Run 'do-release-upgrade' to upgrade to it.

Last login: Sun Mar 23 18:47:53 2025 from 35.235.240.113

```
ana_borrerocabezas@cluster-practica-hadoop-m:~$ sudo sed -i '$d' /etc/hive/conf.dist/hive-site.xml
ana_borrerocabezas@cluster-practica-hadoop-m:~$ sudo sed -i '$a \ <property>\n    <name>es.nodes</name>\n    <value>34.175.180.198</value>\n  </property>\n' /etc/hive/conf.dist/hive-site.xml
ana_borrerocabezas@cluster-practica-hadoop-m:~$ sudo sed -i '$a \ <property>\n    <name>es.port</name>\n    <value>9200</value>\n  </property>\n' /etc/hive/conf.dist/hive-site.xml
ana_borrerocabezas@cluster-practica-hadoop-m:~$ sudo sed -i '$a \ <property>\n    <name>es.nodes.wan.only</name>\n    <value>true</value>\n  </property>\n' /etc/hive/conf.dist/hive-site.xml
ana_borrerocabezas@cluster-practica-hadoop-m:~$ sudo sed -i '$a \ <property>\n    <name>hive.aux.jars.path</name>\n    <value>/usr/lib/hive/lib/elasticsearch-hadoop-8.14.1.jar,/usr/lib/hive/lib/commons-httpclient-3.1.jar</value>\n  </property>\n</configuration>' /etc/hive/conf.dist/hive-site.xml
ana_borrerocabezas@cluster-practica-hadoop-m:~$ sudo cp elasticsearch-hadoop-8.14.1.jar /usr/lib/hive/lib/
ana_borrerocabezas@cluster-practica-hadoop-m:~$ sudo cp commons-httpclient-3.1.jar /usr/lib/hive/lib/
ana_borrerocabezas@cluster-practica-hadoop-m:~$ sudo cat /etc/hive/conf.dist/hive-site.xml
<?xml version="1.0" ?>
<!--
Licensed to the Apache Software Foundation (ASF) under one or more
contributor license agreements.  See the NOTICE file distributed with
this work for additional information regarding copyright ownership.
The ASF licenses this file to You under the Apache License, Version 2.0
(the "License"); you may not use this file except in compliance with
```

PART 2-



SSH en el navegador

↑ SUBIR ARCHIVO

↓ DESCARGAR ARCHIVO



```
<property>
  <name>hive.fetch.task.conversion</name>
  <value>minimal</value>
</property>
<property>
  <name>es.nodes</name>
  <value>34.175.180.198</value>
</property>

<property>
  <name>es.port</name>
  <value>9200</value>
</property>

<property>
  <name>es.nodes.wan.only</name>
  <value>true</value>
</property>

<property>
  <name>hive.aux.jars.path</name>
  <value>/usr/lib/hive/lib/elasticsearch-hadoop-8.14.1.jar,/usr/lib/hive/lib/commons-httpclient-3.1.jar</value>
</property>
</configuration>
ana_borrerocabezas@cluster-practica-hadoop-m:~$ ^C
ana_borrerocabezas@cluster-practica-hadoop-m:~$ ls -lh /usr/lib/hive/lib/ | grep elasticsearch-hadoop
-rw-r--r-- 1 root root 2.2M Mar 23 18:56 elasticsearch-hadoop-8.14.1.jar
ana_borrerocabezas@cluster-practica-hadoop-m:~$ ls -lh /usr/lib/hive/lib/ | grep commons-httpclient
-rw-r--r-- 1 root root 298K Mar 23 18:56 commons-httpclient-3.1.jar
ana_borrerocabezas@cluster-practica-hadoop-m:~$ ^C
ana_borrerocabezas@cluster-practica-hadoop-m:~$ sudo systemctl restart hive-server2
ana_borrerocabezas@cluster-practica-hadoop-m:~$ sudo systemctl status hive-server2
● hive-server2.service - LSB: Hive Server2
   Loaded: loaded (/etc/init.d/hive-server2; generated)
   Drop-In: /etc/systemd/system/hive-server2.service.d
            └─hive-hbase.conf, restart.conf
   Active: active (running) since Sun 2025-03-23 19:02:49 UTC; 13s ago
     Docs: man:systemd-sysv-generator(8)
  Process: 40315 ExecStart=/etc/init.d/hive-server2 start (code=exited, status=0/SUCCESS)
 Main PID: 40334 (java)
    Tasks: 0 (limit: 9517)
   Memory: 4.0K
    CGroup: /system.slice/hive-server2.service
            └─ 40334 /usr/lib/jvm/temurin-11-jdk-amd64/bin/java -Dproc_jar -Dhive.log.dir=/var/log/hive -Dhive.lo

Mar 23 19:02:46 cluster-practica-hadoop-m systemd[1]: Starting LSB: Hive Server2...
Mar 23 19:02:46 cluster-practica-hadoop-m su[40332]: (to hive) root on none
Mar 23 19:02:46 cluster-practica-hadoop-m su[40332]: pam_unix(su:session): session opened for user hive(uid=119)
Mar 23 19:02:46 cluster-practica-hadoop-m su[40332]: pam_unix(su:session): session closed for user hive
Mar 23 19:02:49 cluster-practica-hadoop-m hive-server2[40315]: * Started Hive Server2 (hive-server2):
Mar 23 19:02:49 cluster-practica-hadoop-m systemd[1]: Started LSB: Hive Server2.
lines 1-19/19 (END)
```